

Streams - CPS



efficient functional style
pipelines on streams of data
{no only in F#}

“Although threads seem to be a small step from sequential computation, in fact, they represent a huge step. They discard the most essential and appealing properties of sequential computation: understandability, predictability, and determinism. Threads, as a model of computation, are wildly nondeterministic, and the job of the programmer becomes one of pruning that non-determinism.”

— Edward A. Lee

(*The Problem with Threads, Berkeley 2006*)

AGENDA

Different Streams Implementation, *optimization* and benchmarking

Lighting fast Java 8 Streams – How is it possible?

Continuation Passing Style (CPS)

Create your own Fast Streams

Demo and performance measurements

SOMETHING ABOUT ME

RICCARDO TERRELL

- Originally from Italy, currently living in USA ~10 years
 - Living/working in Washington DC
- +/- 18 years in professional programming
 - C++/VB → Java → .Net C# → Haskell → C# & F# → Scala → ??
- Organizer of the DC F# User Group
- Software Architect @ Microsoft
- Believer in polyglot programming as a mechanism in finding the right tool for the job



@trikace

rickyterrell.com



GOALS - PLAN OF THE TALK

Streams are
lightweight and
powerful

CPS is a great,
useful but miss-
used tool
- to extend functionality

There is always
room to improve
performance



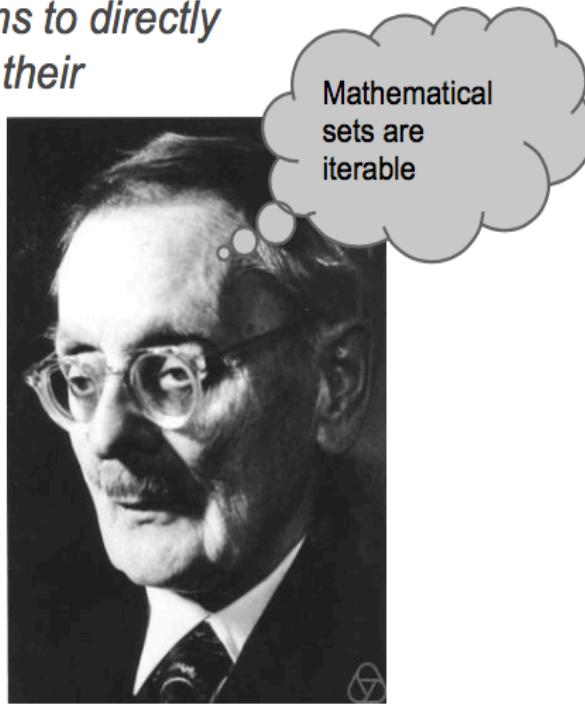


Streams

do not provide a means to directly access or manipulate their elements

```
interface Iterable<T> {  
    Iterator<T> iterator()  
}
```

```
interface Iterator<T> {  
    boolean hasNext()  
    T next()  
}
```



Ernst Friedrich Ferdinand Zermelo (1871–1953)

http://en.wikipedia.org/wiki/Ernst_Zermelo#mediaviewer/File:Ernst_Zermelo.jpeg

Typical Pipeline Pattern



***list comprehension is a syntactic construct
for creating a list based on existing lists***

***It follows the form of the mathematical
“set comprehension” to use HOF as map / filter / reduce ...***



Typical Pipeline Pattern

```
// source |>  
// (('a -> bool) -> 'a list -> 'a list)  
// (('a -> 'b) -> 'a list -> 'b list)  
// (('a -> 'a -> 'a) -> 'a list -> 'a)  
let source' = source  
    |> List.filter f |> List.map m |> List.reduce(+)
```

- inter : intermediate lazy operations (map, filter...)
- terminal : produces result or side-effects (reduce...)



JAVA 8 -> STREAM

```
public double sumOfSquaresSeq(double[] v) {  
    double sum = DoubleStream.of(v)  
        .map(d -> d * d)  
        .sum(); // (x,y) -> x + y  
    return sum;  
}
```



JAVA 8 -> STREAMS -> PARALLEL

```
public double sumOfSquaresSeq(double[] v) {  
    double sum = DoubleStream.of(v)  
        .parallel()  
        .map(d -> d * d)  
        .sum(); // (x,y) -> x + y  
  
    return sum;  
}
```



JAVA 8 -> STREAMS -> COMPILER

```
<T, R> Stream<R> map(Stream<T> source, Function<T, R> mapper) {  
    return new MapperStream<T, R>(source) {  
        Consumer<T> wrap(Consumer<R> consumer) {  
            return new Consumer<T>() {  
                void accept(T v) {  
                    consumer.accept(mapper.apply(v));  
                }  
            };  
        }  
    };  
}
```



JAVA 8 -> STREAMS -> COMPILER

```
<T, R> Stream<R> map(Stream<T> source, Function<T, R> mapper) {  
    return new MapperStream<T, R>(source) {  
        Consumer<T> wrap(Consumer<R> consumer) {  
            return new Consumer<T>() {  
                void accept(T v) {  
                    consumer.accept(mapper.apply(v));  
                }  
            };  
        }  
        do { consumer.accept(a[i]); } while (++i < hi);  
    }  
}
```



C# -> LINQ

```
public int sumOfSquaresEvenSeq(int[] v) {  
    int sum =  
        v.Where(x => x % 2 == 0)  
            .Select(x => x * x)  
            .Sum();  
  
    return sum;  
}
```



C# -> LINQ

// Fluent Style

```
nums.Where(x => x % 2 == 0).Select(x => x * x).Sum();
```

// Query Syntax

```
(from x in nums where x % 2 == 0 select x * x).Sum();
```



C# -> LINQ -> PARALLEL

```
public int sumOfSquaresEvenSeq(int[] v) {  
  
    int sum =  
  
        v.AsParallel()  
            .Where(x => x % 2 == 0)  
            .Select(x => x * x)  
            .Sum();  
  
    return sum;  
}
```



F# -> Seq

```
let data = [|1..1000000|] |> Array.map int64  
  
data  
|> Seq.filter(fun i -> i % 2L = 0L) // lazy  
|> Seq.map(fun i -> i + 1L) // lazy  
|> Seq.sum // eager force evaluation  
  
|> ( 'a -> ('a -> 'b) -> 'b)
```



F# -> Seq -> Parallel

```
let data = [|1..1000000|] |> Array.map int64  
  
data  
|> PSeq.filter(fun i -> i % 2L = 0L) // lazy  
|> PSeq.map(fun i -> i + 1L) // lazy  
|> PSeq.sum // eager force evaluation
```



C# & F# -> LINQ & SEQ -> COMPILER

```
Interface IEnumerable<T> { IEnumerator<T> GetEnumerator(); }
```

```
Interface IEnumerator<T> { // Return current position element  
    T Current { get; } // Move to next element,  
                      // returns false if no more elements  
    bool MoveNext();  
}
```



C# & F# -> LINQ & Seq -> Compiler

```
class SelectEnumerator<T, R> : IEnumerator<R> {  
    private readonly IEnumerator<T> inner;  
    private readonly Func<T, R> func;  
    public SelectEnumerator(IEnumerator<T> inner, Func<T, R> func) {  
        this.func = func;  
        this.inner = inner;  
    }  
    MoveNext() { return inner.MoveNext(); }  
    Current { get { return func(inner.Current); } }  
}
```



C# & F# -> LINQ & Seq -> Compiler

```
class SelectEnumerator<T, R> : Ienumerator<R> {  
    private readonly Ienumerator<T> inner;  
    private readonly Func<T, R> func;  
    public SelectEnumerator(Ienumerator<T> inner, Func<T, R> func) {  
        this.func = func;  
        this.inner = inner;  
    }  
    MoveNext() { return inner.MoveNext(); }  
    Current { get { return func(inner.Current); } }  
}
```



SCALA -> SEQVIEW

```
def sumOfSqaureSeq(v: Array[Double]) : Double =  
    val sum : Double = v  
        .view  
        .map(d => d * d)  
        .sum  
    sum  
}
```



SCALA -> SEQVIEW -> COMPILER

```
def map[T, U](source: Iterable[T], f: T => U) =  
    new Iterable[U] {  
        def iterator = source.iterator map f  
    }
```

```
def map[T, U](source: Iterator[T], f: T => U): Iterator[U] =  
    new Iterator[U] {  
        def hasNext = source.hasNext  
        def next() = f(source.next())  
    }
```



SCALA -> SEQVIEW -> COMPILER

```
def map[T, U](source: Iterable[T], f: T => U) =  
    new Iterable[U] {  
        def iterator = source.iterator map f  
    }
```

```
def map[T, U](source: Iterator[T], f: T => U): Iterator[U] =  
    new Iterator[U] {  
        def hasNext = source.hasNext  
        def next() = f(source.next())  
    }
```



Clash of the Lambdas

Clash of the Lambdas

Through the Lens of Streaming APIs

Aggelos Biboudis

University of Athens
biboudis@di.uoa.gr

Nick Palladinos

Nessos Information Technologies, SA
npal@nessos.gr

Yannis Smaragdakis

University of Athens
smaragd@di.uoa.gr

Abstract

The introduction of lambdas in Java 8 completes the slate of statically-typed, mainstream languages with both object-oriented and functional features. The main motivation for lambdas in Java has been to facilitate stream-based declarative APIs, and, therefore, easier parallelism. In this paper, we evaluate the performance impact of lambda abstraction employed in stream processing, for a variety of high-level languages that run on a virtual machine (C#, F#, Java and Scala) and runtime platforms (JVM on Linux and Windows, .NET CLR for Windows, Mono for Linux). Furthermore, we evaluate the performance gain that two optimizing libraries (ScalaBlitz and LinqOptimizer) can offer for C#, F# and Scala. Our study is based on small-scale throughput-benchmarking, with significant care to isolate different factors, consult experts on the sys-

in lines 3,4 use only their argument values, and no values from the environment.

```
1 public int sumOfSquaresEvenSeq(int[] v) {
2     int sum = IntStream.of(v)
3         .filter(x -> x % 2 == 0)
4         .map(x -> x * x)
5         .sum();
6     return sum;
7 }
```

The above computation can be trivially parallelized with the addition of a `.parallel()` operator before the call to `filter`. This ability showcases the simplicity benefits of streaming abstractions for parallel operations.

The background of the slide features a dark, abstract image of blurred light streaks, resembling a road at night or data flow, which serves as a visual metaphor for speed and data processing.

Make functional data query pipelines **FAST**



- **LinqOptimizer**
- **Steno**
- **ScalaBlitz**
- **Stream Fusion**



An automatic query optimizer for the compiler that target Sequences

0010001001010101010102001

- Query optimizers into fast loop-based imperative code
- Increase performance of up to ~15x



Linq-Optimizer

```
var query = (from num in nums.AsQueryable ()  
            where num % 2 == 0  
            select num * num).Sum ();
```



Linq-Optimizer

```
var query = (from num in nums.AsQueryable ()  
            where num % 2 == 0  
            select num * num).Sum ();
```

Compiles
to

```
int sum = 0;  
for (int index = 0; index < nums.Count; index++) {  
    int num = nums [index];  
    if (num % 2 == 0)  
        sum += num * num;  
}
```



Loop Fusion

```
int i, a[100], b[100];
for (i = 0; i < 100; i++)
    a[i] = 1;
for (i = 0; i < 100; i++)
    b[i] = 2;
```

Is equivalent to

```
int i, a[100], b[100];
for (i = 0; i < 100; i++)
{
    a[i] = 1;
    b[i] = 2;
}
```



Linq-Optimizer

```
var query = (from num in nums.AsParallelQueryExpr()
            where num % 2 == 0
            select num * num).Sum();
```



Linq-Optimizer

```
var query = (from num in nums. nums.AsParallelQueryExpr()
             where num % 2 == 0
             select num * num).Sum();
```

Compiles
to

```
ParallelQuery.ReduceCombine (nums, 0,
    (acc, num) => {
        if (num % 2 == 0)
            return acc + num * num;
        else
            return acc;
    }, (left, right) => left + right);
```



Linq-Optimizer disadvantages

- Runtime compilation
 - Overhead (mitigated by caching)
 - Emitting IL not cross-platform
 - Security restriction in cloud, mobile
 - Access to private fields/methods?
 - Compiler changes ?



Steno: Automatic Optimization of Declarative Queries

Derek G. Murray

University of Cambridge Computer Laboratory
Derek.Murray@cl.cam.ac.uk

Michael Isard Yuan Yu

Microsoft Research Silicon Valley
{misard, yuanbyu}@microsoft.com

Abstract

Declarative queries enable programmers to write data manipulation code without being aware of the underlying data structure implementation. By increasing the level of abstraction over imperative code, they improve program readability and, crucially, create opportunities for automatic parallelization and optimization. For example, the Language Integrated Query (LINQ) extensions to C# allow the same declarative query to process in-memory collections, and datasets that are distributed across a compute cluster. However, our experiments show that the *serial* performance of declarative code is several times slower than the equivalent hand-optimized code, because it is implemented using run-time abstractions—such as iterators—that incur overhead due to virtual function calls and

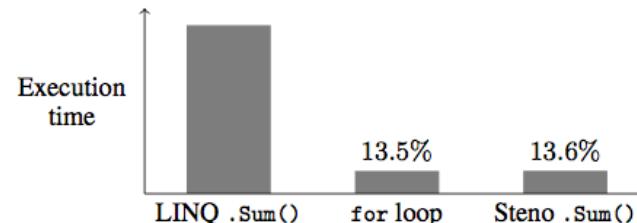


Figure 1. Relative execution time for computing the sum of squares of 10^7 doubles using LINQ, an imperative loop, and a Steno-optimized query. Steno achieves a 7.4× speedup over LINQ.

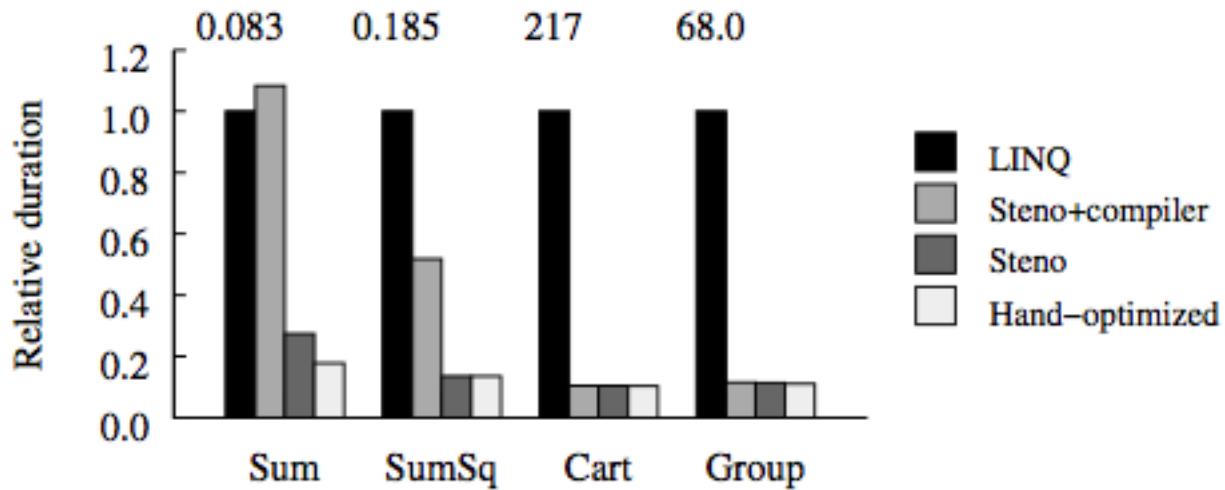


C# & F# LINQ -> STENO

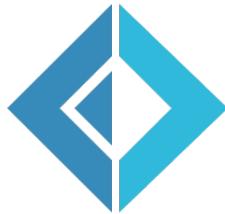
- LINQ are lazy evaluated by using iterator, which is the cause of two virtual function calls per element per query operator
- Nested LINQ query increase overhead previous point
- Steno converts queries into efficient, loop-based imperative code



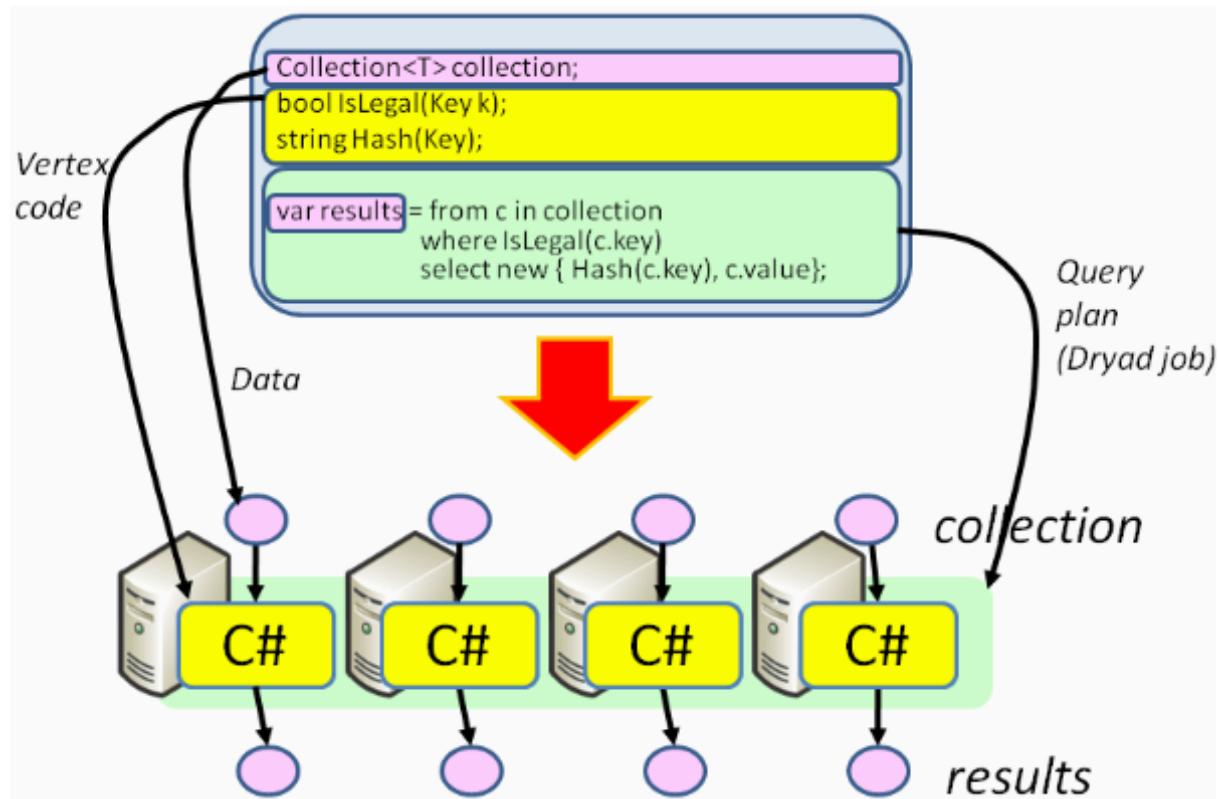
C# & F# LINQ -> STENO



Relative performance of LINQ, Steno-optimized and hand-optimized queries for sequential micro benchmarks. Lower values are better. Each query is annotated with the absolute LINQ execution time, in seconds.



C# & F# LINQ -> STENO -> DRYADLINQ



HASKELL -> STREAM FUSION



Stream Fusion

From Lists to Streams to Nothing at All

Duncan Coutts¹ Roman Leshchinskiy² Don Stewart²

¹ Programming Tools Group
Oxford University Computing Laboratory
duncan.coutts@comlab.ox.ac.uk

² Computer Science & Engineering
University of New South Wales
{rl,dons}@cse.unsw.edu.au

Abstract

This paper presents an automatic deforestation system, *stream fusion*, based on equational transformations, that fuses a wider range of functions than existing short-cut fusion systems. In particular, stream fusion is able to fuse zips, left folds and functions over nested lists, including list comprehensions. A distinguishing feature of the framework is its simplicity: by transforming list functions to expose their structure, intermediate values are eliminated by general purpose compiler optimisations.

No previously implemented short-cut fusion system eliminates all the lists in this example. The fusion system presented in this paper does. With this system, the Glasgow Haskell Compiler (The GHC Team 2007) applies all the fusion transformations and is able to generate an efficient “worker” function f' that uses only unboxed integers ($\text{Int}\#$) and runs in constant space:

```
f' :: Int# → Int#
f' n =
let go :: Int# → Int# → Int#
in
```



HASKELL -> STREAM FUSION

Stream Fusion

From Lists to Streams to Nothing at All

Duncan Coutts¹ Roman Leshchinskiy² Don Stewart²

¹ Programming Tools Group
Oxford University Computing Laboratory
duncan.coutts@comlab.ox.ac.uk

² Computer Science & Engineering
University of New South Wales
{rl,dons}@cse.unsw.edu.au

The key trick is that all producers are non-recursive

Abstract

This paper presents an automatic deforestation system, *stream fusion*, based on equational transformations, that fuses a wider range of functions than existing short-cut fusion systems. In particular, stream fusion is able to fuse zips, left folds and functions over nested lists, including list comprehensions. A distinguishing feature of the framework is its simplicity: by transforming list functions to expose their structure, intermediate values are eliminated by general purpose compiler optimisations.

No previously implemented short-cut fusion system eliminates all the lists in this example. The fusion system presented in this paper does. With this system, the Glasgow Haskell Compiler (The GHC Team 2007) applies all the fusion transformations and is able to generate an efficient “worker” function f' that uses only unboxed integers ($\text{Int}\#$) and runs in constant space:

```
f' :: Int# → Int#
f' n =
  let go :: Int# → Int# → Int#
    in ...
```



Haskell -> STREAM FUSION

```
data Stream a = E s . Stream (s → Step a s) s

maps :: (a → b) → Stream a → Stream b
maps f (Stream next0 s0) = Stream next s0
  where
    next s = case next0 s of
      Done → Done
      Skip s' → Skip s'
      Yield x s' → Yield (f x) s'
```

HASKELL -> STREAM FUSION



Stream fusion three step approach:

1. Convert recursive structures into non-recursive *co-structures*
2. Eliminate superfluous conversions between structures and co-Structures
3. Use general optimizations to fuse the co-structure code.

HASKELL -> STREAM FUSION



Haskell Beats C Using Generalized Stream Fusion

Geoffrey Mainland

Microsoft Research Ltd
Cambridge, England
gmainlan@microsoft.com

Roman Leshchinskiy

rl@cse.unsw.edu.au

Simon Peyton Jones

Microsoft Research Ltd
Cambridge, England
simonpj@microsoft.com

Abstract

Stream fusion [6] is a powerful technique for automatically transforming high-level sequence-processing functions into efficient implementations. It has been used to great effect in Haskell libraries for manipulating byte arrays, Unicode text, and unboxed vectors. However, some operations, like vector append, still do not perform well within the standard stream fusion framework. Others, like SIMD computation using the SSE and AVX instructions available on modern x86 chips, do not seem to fit in the framework at all.

which fusion is relatively straightforward, was the first truly general solution.

Instead of working directly with lists or vectors, stream fusion works by re-expressing these functions as operations over streams, each represented as a state and a step function that transforms the state while potentially yielding a single value. Alas, different operations need different stream representations, and no single representation works well for all operations (§2.3). Furthermore, for many operations it is not obvious what the choice of representation should be.

HASKELL -> STREAM FUSION



Haskell Beats C Using Generalized Stream Fusion

Geoffrey Mainland

Microsoft Research Ltd
Cambridge, England
gmainlan@microsoft.com

Roman Leshchinskiy

rl@cse.unsw.edu.au

Simon Peyton Jones

Microsoft Research Ltd
Cambridge, England
simonpj@microsoft.com

Haskell – Stream Fusion faster than C

Abstract

Stream fusion [6] is a powerful technique for automatically transforming high-level sequence-processing functions into efficient implementations. It has been used to great effect in Haskell libraries for manipulating byte arrays, Unicode text, and unboxed vectors. However, some operations, like vector append, still do not perform well within the standard stream fusion framework. Others, like SIMD computation using the SSE and AVX instructions available on modern x86 chips, do not seem to fit in the framework at all.

which fusion is relatively straightforward, was the first truly general solution.

Instead of working directly with lists or vectors, stream fusion works by re-expressing these functions as operations over streams, each represented as a state and a step function that transforms the state while potentially yielding a single value. Alas, different operations need different stream representations, and no single representation works well for all operations (§2.3). Furthermore, for many operations it is not obvious what the choice of representation should be.



ScalaBlitz

Lightning-fast Scala collections framework



SCALABLITZ

```
def sumOfSquareSeqBlitz (v : Array[Double]) : Double = {  
    optimize {  
        val sum : Double = v  
            .map(d => d * d)  
            .sum  
        sum  
    }  
}
```

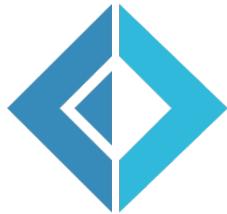
```
def optimize[T](exp: T): Any = macro optimize_impl[T]
```



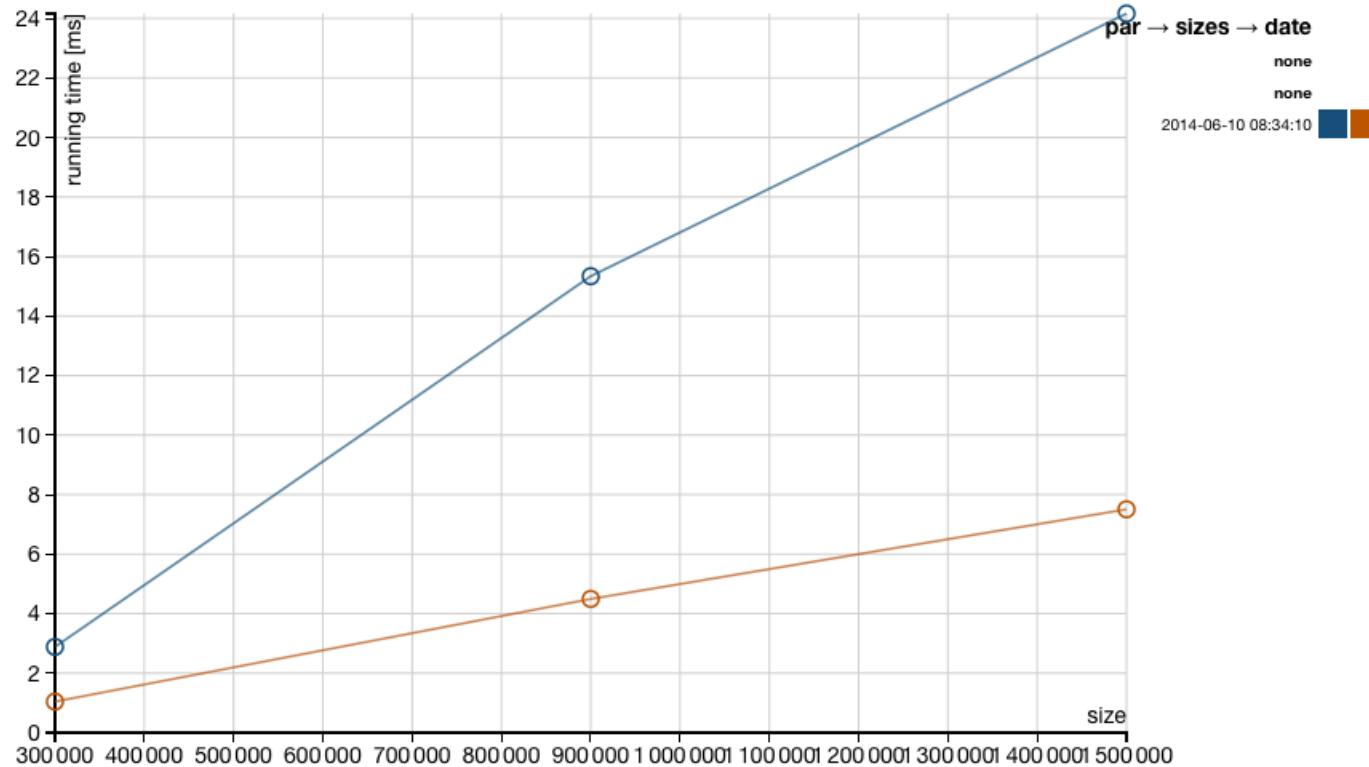
SCALABLITZ

```
def sumOfSquareSeaBlitz (v : Array[Double]) : Double = {  
    optimize {  
        val sum : Double = v  
        .map(x => x * x)  
        .sum  
    }  
}
```

```
def optimize[T](exp: T): Any = macro optimize_impl[T]
```



SCALBLITZ - REDUCE





Micro-benchmark results

Java vs Scala vs C# & F#

Sequential & Parallel

- Sum iteration speed with no lambdas, just a single iteration.
- Sum Of Squares a small pipeline with one map operation
- Cartesian product a nested pipeline with a flatMap and an inner operation, again with a flatMap (capturing a variable), to encode a Cartesian product.

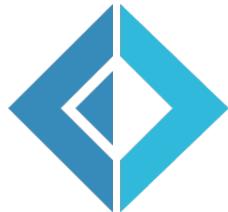


Settings & Setup

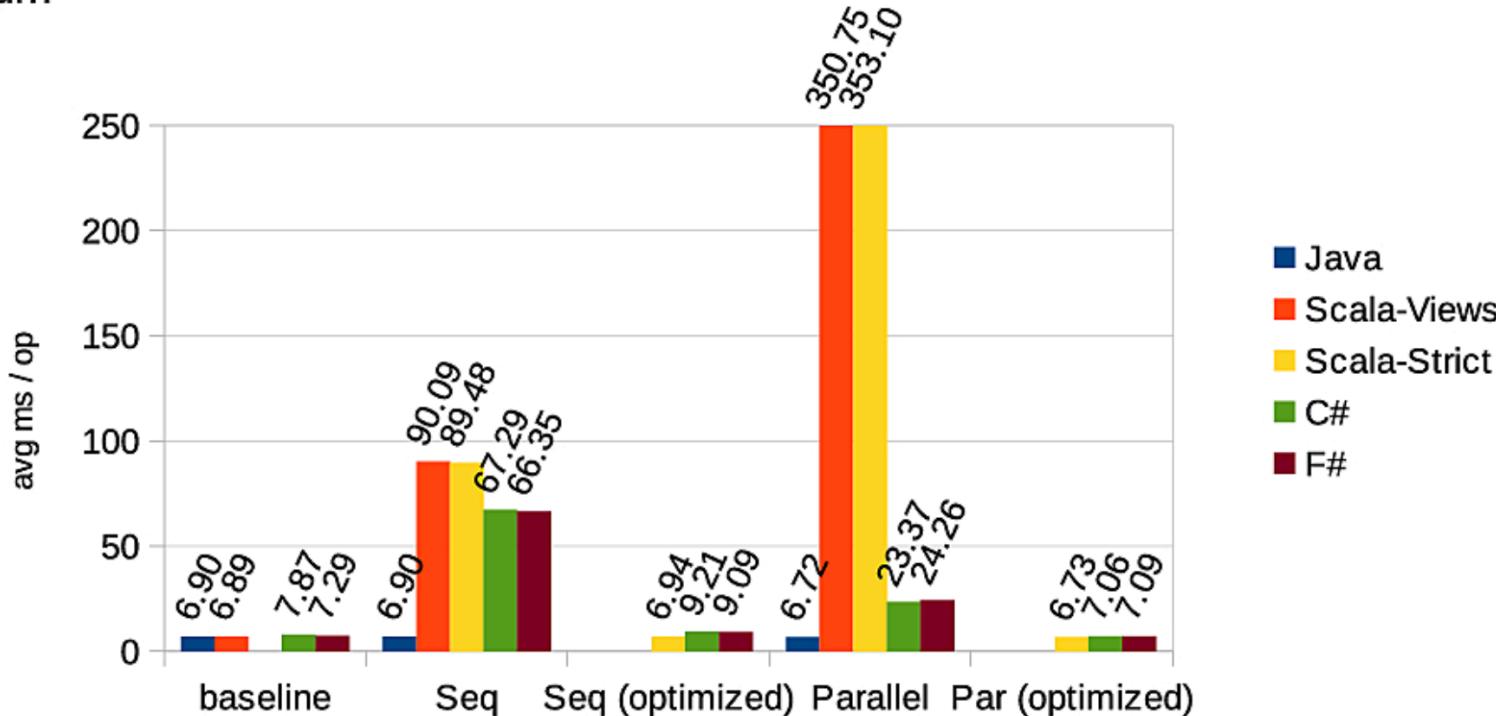
	Windows	Ubuntu Linux
Version	8.1	13.10/3.11.0-20
Architecture	x64	x64
CPU	Intel Core i5-3360M vPro 2.8GHz	
Cores		2 physical x 2 logical
Memory		4GB

	Windows	Ubuntu Linux
Java		Java 8 (b132)/JVM 1.8
Scala		2.10.4/JVM 1.8
C#	C#5 /CLR v4.0	C# mono 3.4.0.0/mono 3.4.0
F#	F#3.1/CLR v4.0	F# open-source 3.0/mono 3.4.0

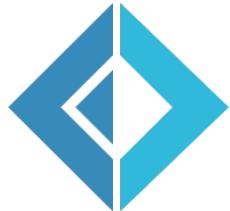
Windows



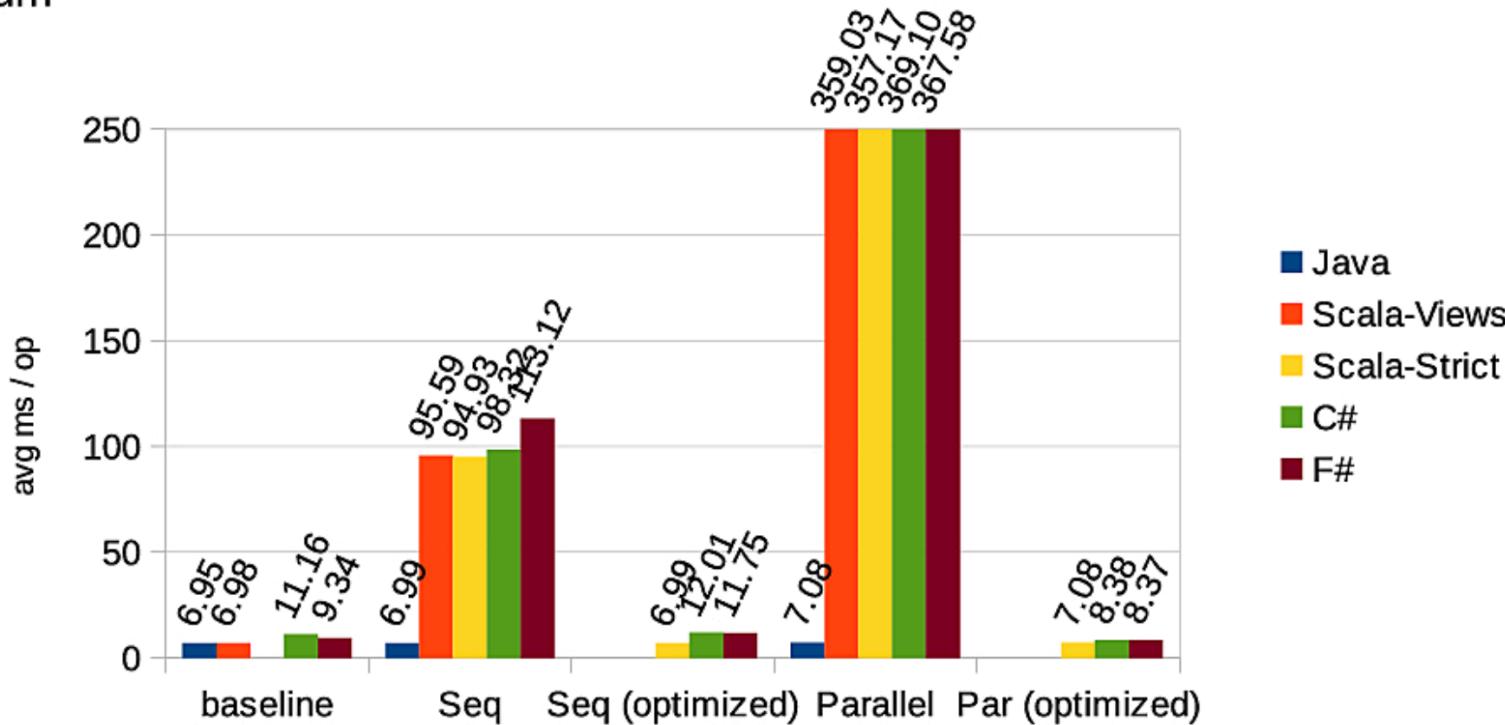
Sum



Linux



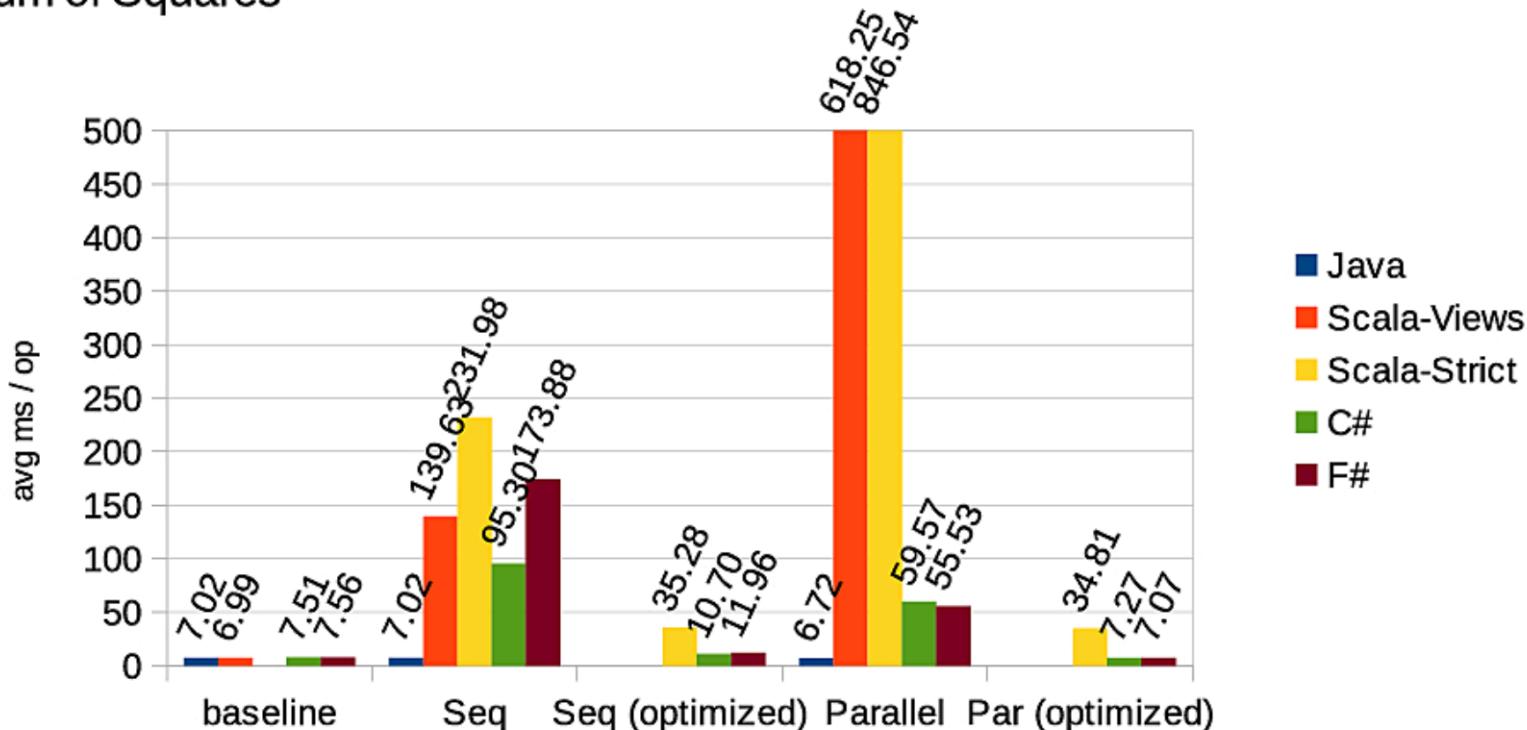
Sum

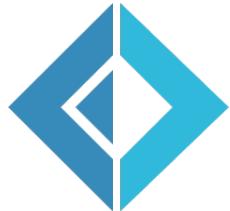




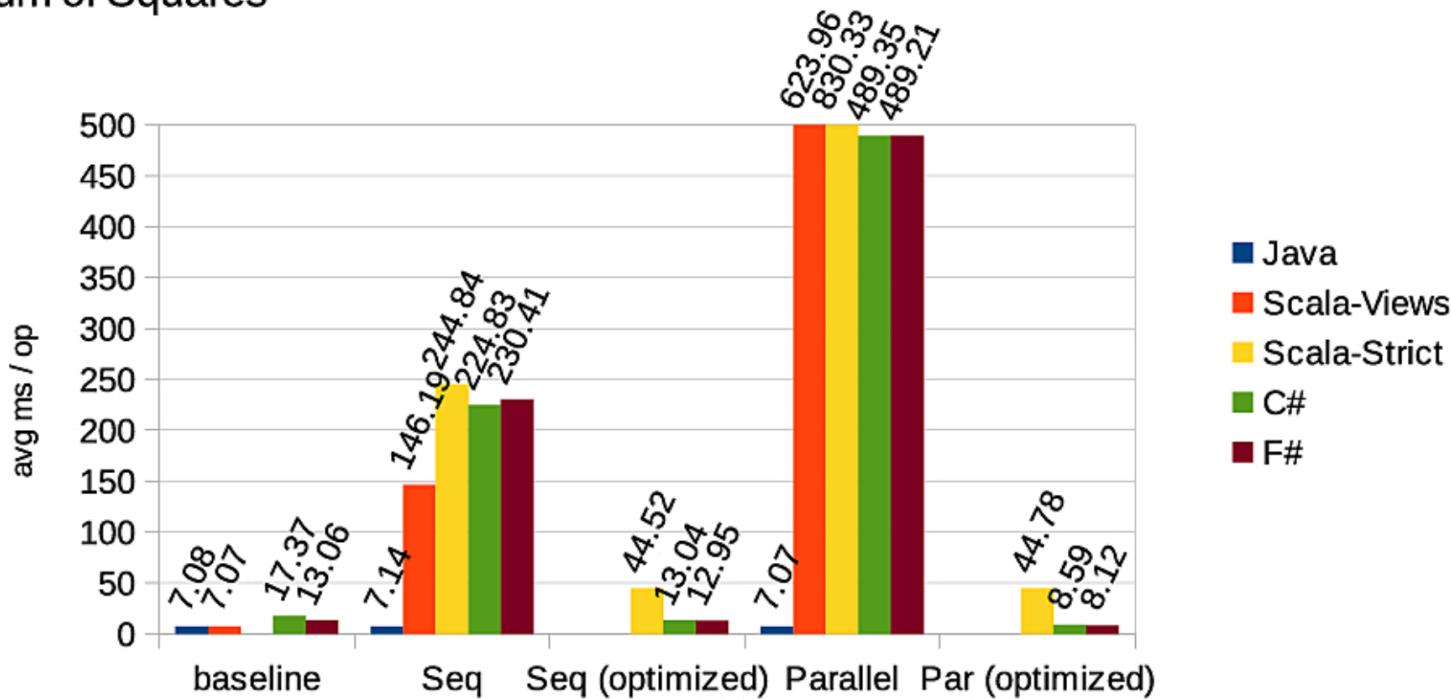
Windows

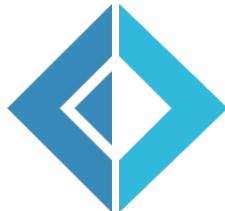
Sum of Squares





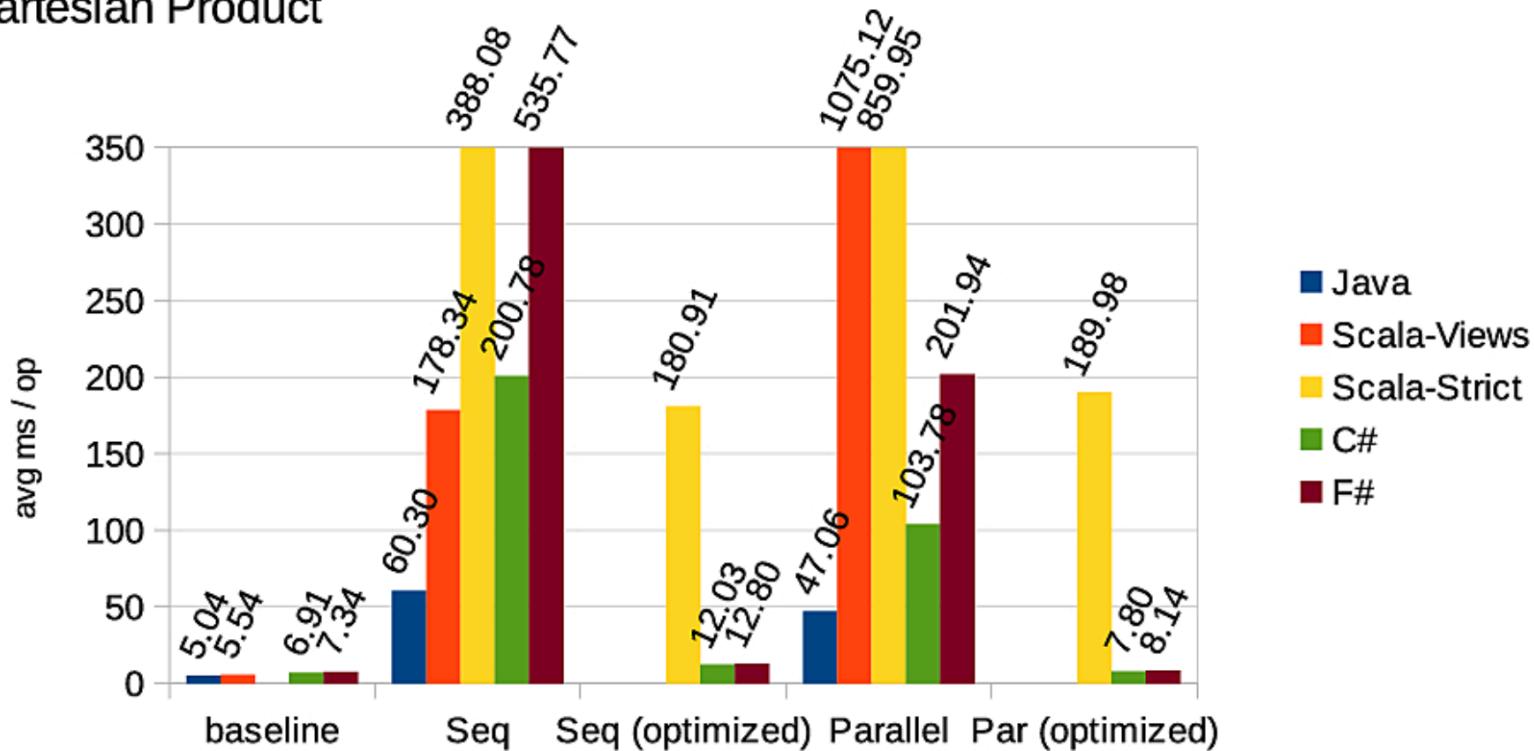
Sum of Squares





Windows

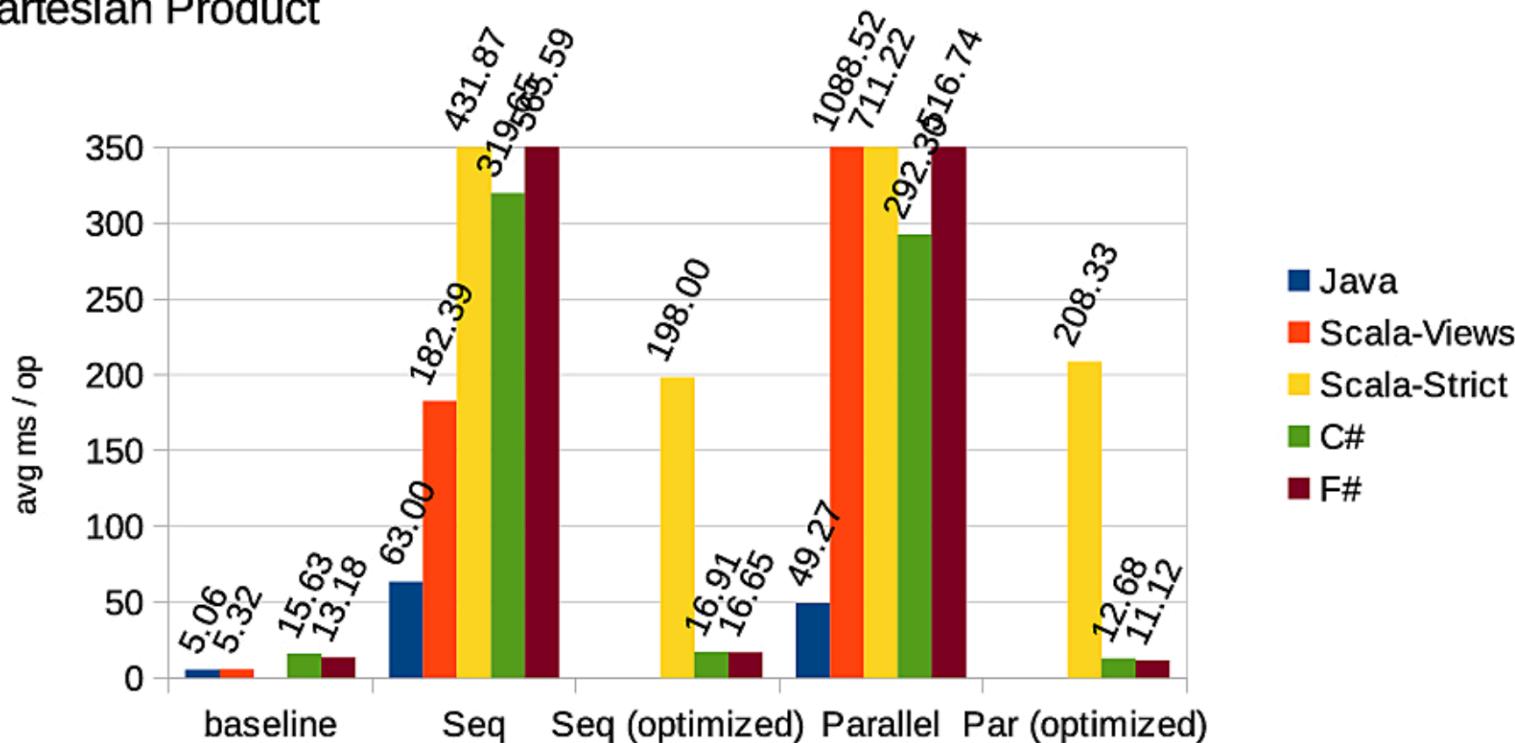
Cartesian Product





Linux

Cartesian Product





Java 8 is very fast



MASSACHUSETTS INSTITUTE OF TECHNOLOGY
ARTIFICIAL INTELLIGENCE LABORATORY

AI Memo No. 349

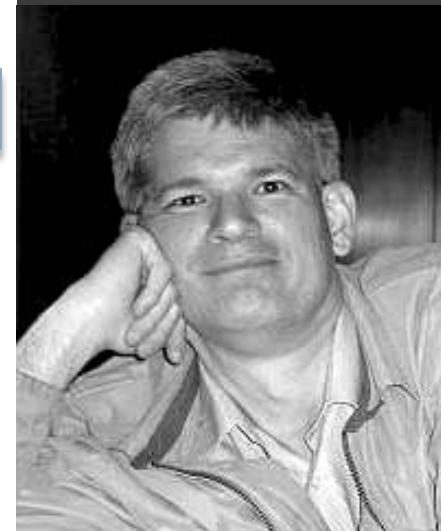
December 1975

SCHEME

AN INTERPRETER FOR EXTENDED LAMBDA CALCULUS

by

Gerald Jay Sussman and Guy Lewis Steele Jr.



Abstract:

Inspired by ACTORS [Greif and Hewitt] [Smith and Hewitt], we have implemented an interpreter for a LISP-like language, SCHEME, based on the lambda calculus [Church], but extended for side effects, multiprocessing, and process synchronization. The purpose of this implementation is tutorial. We wish to:

Another Way To Do Recursion

Now consider the following alternative definition of FACT. It has an extra argument, the continuation [Reynolds], which is a function to call with the answer, when we have it, rather than return a value; that is, rather than ultimately reducing to the desired value, it reduces to a combination which is the application of the continuation to the desired value.



A function written in continuation-passing style takes an extra argument: an explicit "continuation" i.e. a function of one argument. When the CPS function has computed its result value, it "returns" it by calling the continuation function with this value as the argument. That means that when invoking a CPS function, the calling function is required to supply a procedure to be invoked with the subroutine's "return" value. Expressing code in this form makes a number of things explicit which are implicit in direct style. These include: procedure returns, which become apparent as calls to a continuation; intermediate values, which are all given names; order of argument evaluation, which is made explicit; and tail calls, which simply call a procedure with the same continuation, unmodified, that was passed to the caller.



Programs can be automatically transformed from direct style to CPS.


$$(A \rightarrow B) \rightarrow B$$
$$(A, B \rightarrow \text{unit}) \rightarrow \text{unit}$$



A little more about CPS

- Converting to CPS

```
static int Max(int n, int m){  
    return n > m ? n : m;  
}
```



A little more about CPS

- **Converting to CPS**

- Change the return type to void
- Add an extra argument of type `Action<T>` where T is the original return type
- Replace all return statements with invocations of the new continuation argument passing the expression used in the return statement

```
static int Max(int n, int m){  
    return n > m ? n : m;  
}
```

```
static void Max(int n, int m, Action<int> k){  
    var max = n > m ? n : m;  
    k (max);  
}
```



Recursive CPS

```
static void Factorial(int n, Action<int> k){  
    if (n == 0)  
        k (1);  
    else  
        Factorial (n - 1, x => k (n * x));  
}  
  
Factorial(5, x => Console.WriteLine(x));
```



A little more about CPS

- Returning Values with Continuations

```
// ((T -> T) -> T) -> T
T CallCC<T>(Func<Func<T, T>, T> f)
```

```
static int Foo(int n){
    return CallCC<int>(Return => {
        Return(n);
        return n + 1;
    });
}
```

```
Foo(42);
```

Why CPS?



What exactly can a programmer do with CPS besides showoff at parties?

- Compilers use a more thorough CPS transformation to produce an intermediate form amenable to many analysis.
- UI frameworks use CPS to keep the UI responsive while allowing nonlinear program interaction.
- Web servers use CPS to allow computation to flow asynchronously across pages (Callbacks)



Java 8 Streams use CPS -> are pushing

```
Stream.OfArray(data)
    .filter(i -> i % 2 == 0)
    .map(i -> i * i)
    .sum();
```

The source is pushing data down the pipeline



C# & F# -> LINQ & Seq are pulling

```
let data = [|1..1000000|] |> Array.map int64  
  
data  
|> Seq.find(fun i -> i % 2L = 0L) // lazy  
|> Seq.map(fun i -> i + 1L) // lazy  
|> Seq.sum // eager force evaluation
```

The terminal function Sum-Aggregate is pulling data from the pipeline via `IEnumerable.Current` and `IEnumerable.MoveNext()`



How does it work applied to Sequences?





How does it work applied to Sequences?

Can we use the same approach?



LET'S MAKE A SIMPLE STREAMS





Starting from Seq.iter

```
// iter : ('T -> unit) -> seq<'T> -> unit
let iter f values =
    for value in values do
        f value
```



Starting from Seq.iter

```
// iter : ('T -> unit) -> seq<'T> -> unit
let iter f values =
    for value in values do
        f value
```

FLIP THE ARGS

seq<'T> -> ('T -> unit) -> unit



Streams !

```
type Stream<'T> = ('T -> unit) -> unit

// iter : seq<'T> -> Stream<'T>
let iter values f =
    for value in values do
        f value
```

Simple Streams



```
type Stream = ('T -> unit) -> unit
```

```
// map : ('T -> 'R) -> Stream<'T> -> Stream<'R>
```

```
let map f stream =
  fun k -> stream (fun v -> k (f v))
```

```
// filter : ('T -> bool) -> Stream<'T> -> Stream<'T>
```

```
let filter f stream =
  fun k -> stream (fun v -> if f v then k v else ())
```

```
// length : Stream<'T> -> int
```

```
let length stream =
  let counter = ref 0
  stream (fun _ -> counter := counter.Value + 1)
  counter.Value
```



Streams are pushing

```
let data = [|1..1000000|] |> Array.map int64

let stream =
    Stream.ofArray data // source
    |> Stream.filter(fun i -> i % 2L = 0)
    |> Stream.map(fun i -> i + 1L)
    |> Stream.sum
```

The source is pushing data down the pipeline

When to stop pushing ?



```
type Stream<'T> = ('T -> unit) -> unit
```

Stopping push required ->

```
Stream.takeWhile : ('T -> bool) ->  
                    Stream<'T> -> Stream<'T>
```

Stopping push



Change

```
type Stream<'T> = ('T -> unit) -> unit
```

to

```
type Stream<'T> = ('T -> bool) -> unit
```

```
// takeWhile : ('T -> bool) -> Stream<'T> -> Stream<'T>
let takeWhile f stream =
  fun k -> stream (fun v -> if f v then k v else false)
```

Streams can push and pull



```
// Provides functions for iteration
type Iterable<'T> = {
    Bulk : unit -> T
    TryAdvance : unit -> bool }
```

ofArray



```
val ofArray : 'T[] -> Stream<'T>
let ofArray values =
    fun k ->
        let bulk () = for value in values do
            k value
        let index = ref -1
        let tryAdvance () =
            incr index;
            if !index < Array.length values then
                (k values.[!index])
                true
            else
                false
        { Bulk = bulk; TryAdvance = tryAdvance }
```

Parallel Streams



```
let data = [| 1..1000000 |] |> Array.map int  
  
data  
|> ParStream.ofArray  
|> ParStream.filter (fun x -> x % 2L = 0L)  
|> ParStream.map (fun x -> x * x)  
|> ParStream.sum
```

ParStream - Parallel Streams



```
type ParStream<'T> = (unit -> ('T -> unit)) -> unit
```

```
val ofArray : 'T[] -> ParStream<'T>
```

```
let ofArray values =
```

```
    fun thunk ->
```

```
        let forks =
```

```
            values
```

```
            |> partitions
```

```
            |> Array.map (fun p -> (p, thunk ()))
```

```
            |> Array.map (fun (p, k) -> fork p k)
```

```
        join forks
```

ParStream - Parallel Streams



```
type ParStream<'T> = (unit -> ('T -> unit)) -> unit

val map : ('T -> 'R) -> ParStream<'T> -> ParStream<'R>
let map f stream =
    fun thunk ->
        stream (fun () ->
            let k = thunk ()
            (fun v -> k (f v)))
```

ParStream - Parallel Streams



```
type ParStream<'T> = (unit -> ('T -> unit)) -> unit
```

```
val sum : ParStream<'T> -> int
```

```
let sum stream =
```

```
    let array = new ResizeArray<int ref>()
```

```
    stream (fun () ->
```

```
        let sum = ref 0
```

```
        array.Add(sum)
```

```
        (fun v -> sum := sum + v)
```

```
)
```

```
array |> Array.map (fun sum -> !sum) |> Array.sum
```

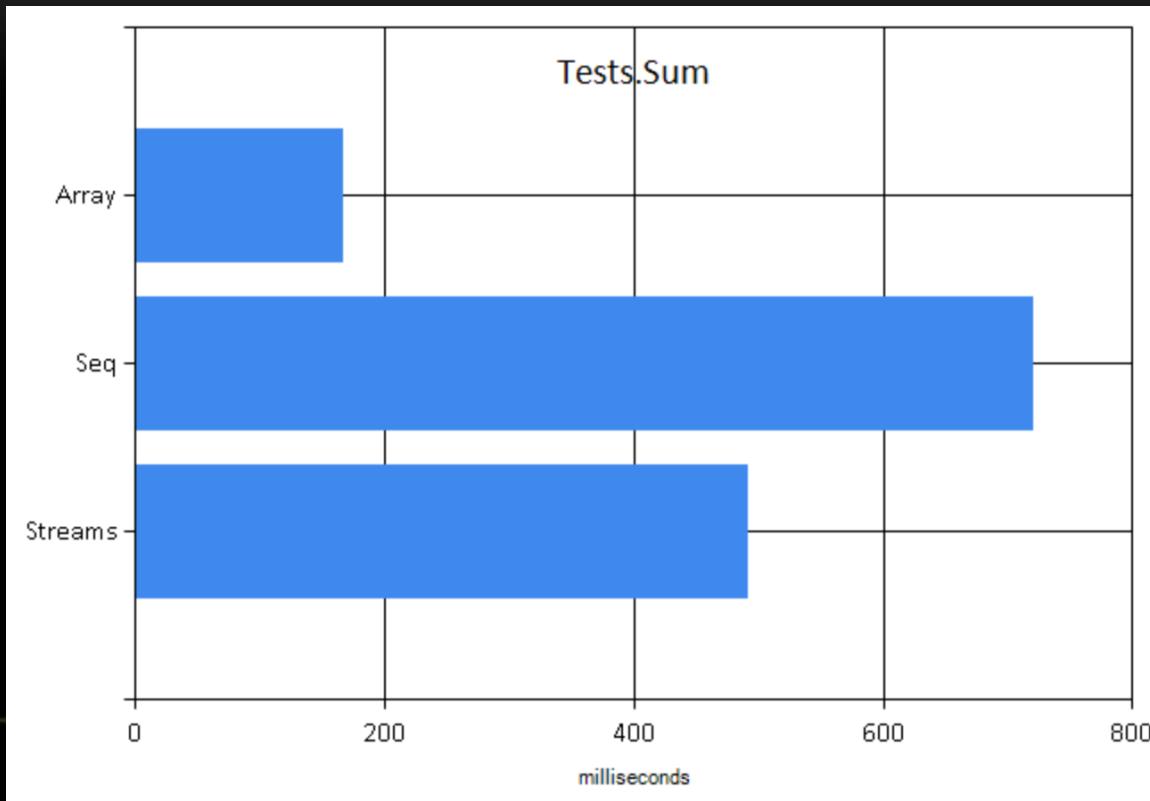


Benchmarks

i7 x 3.7 Ghz (4 physical), 6 GB RAM

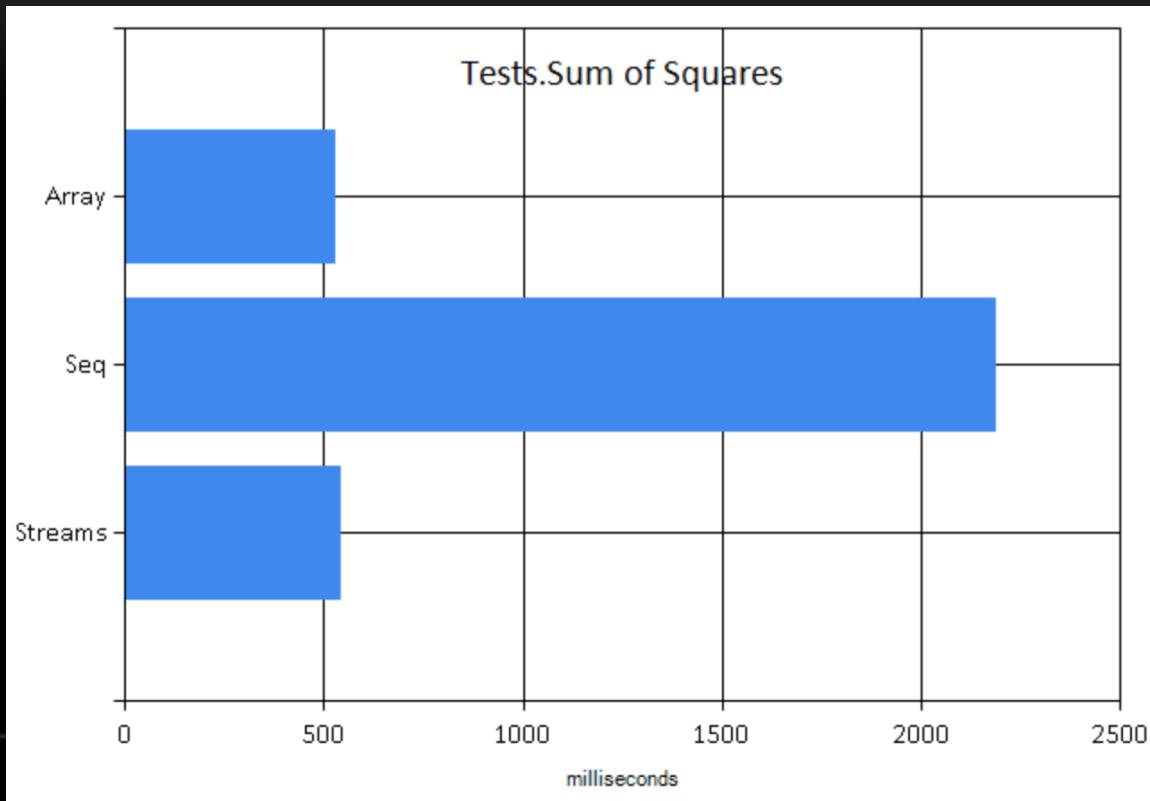


Sum



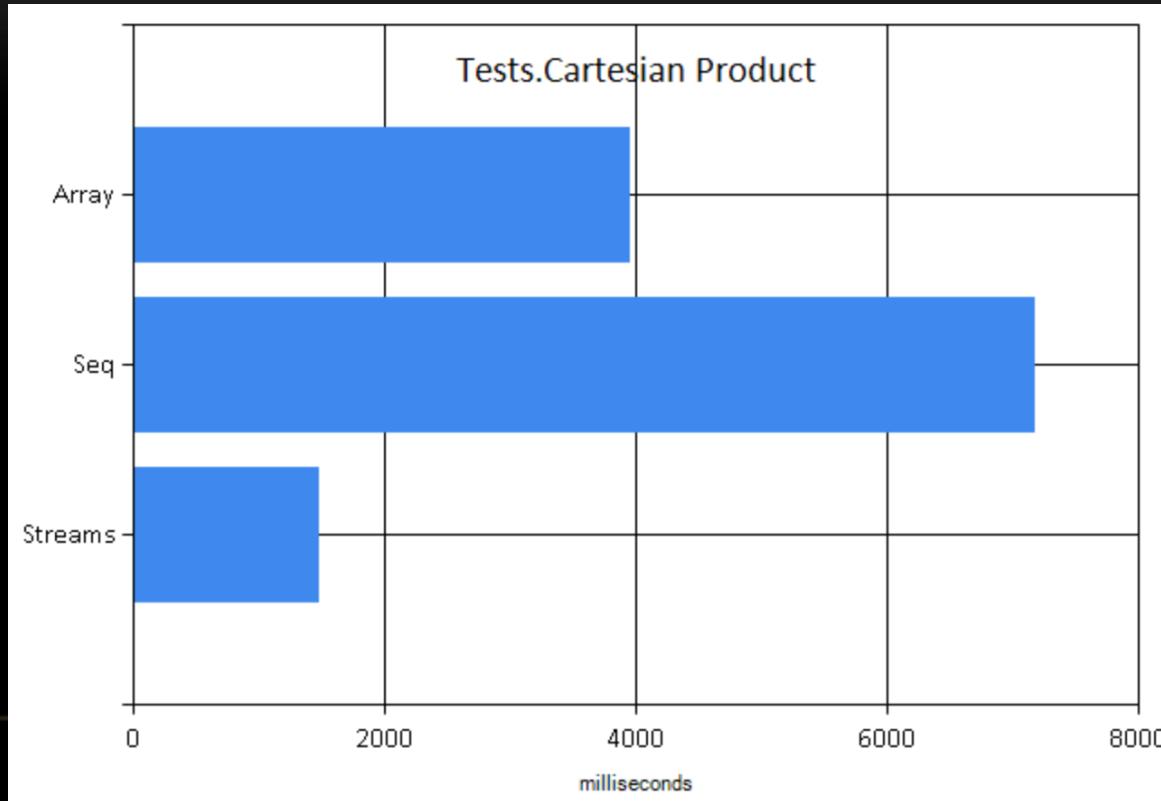


Sum of squares



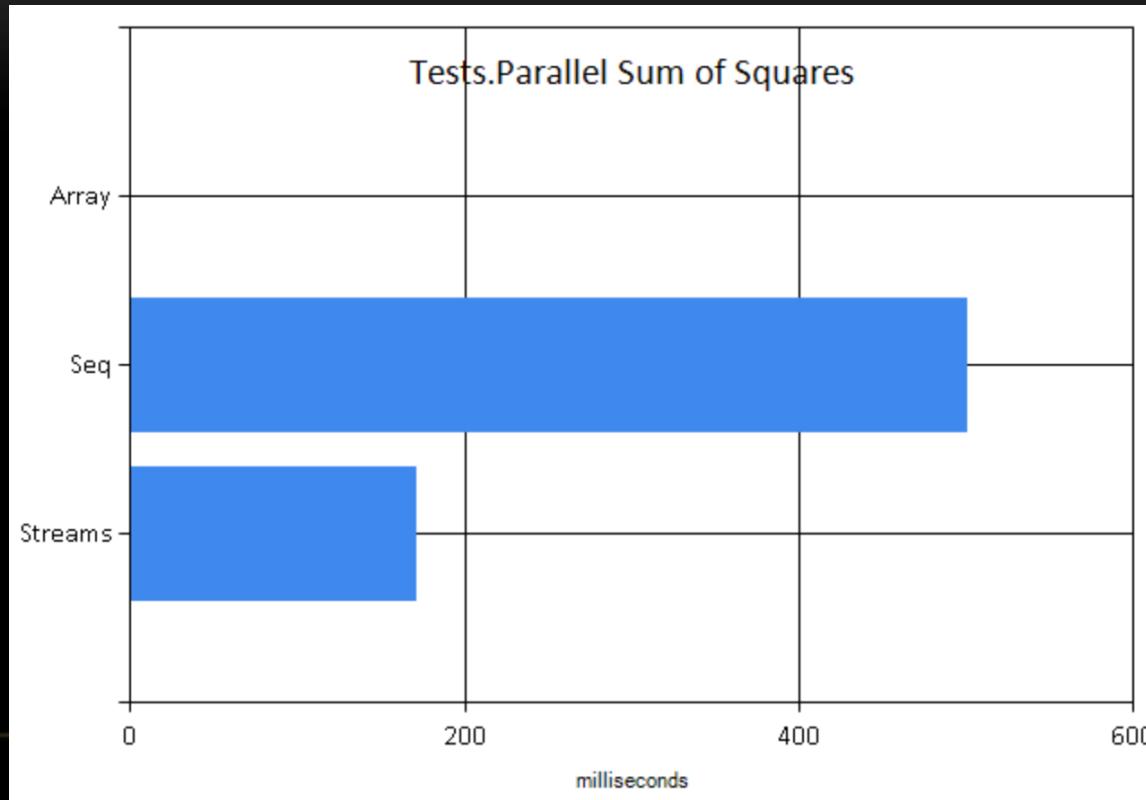


Cartesian Product

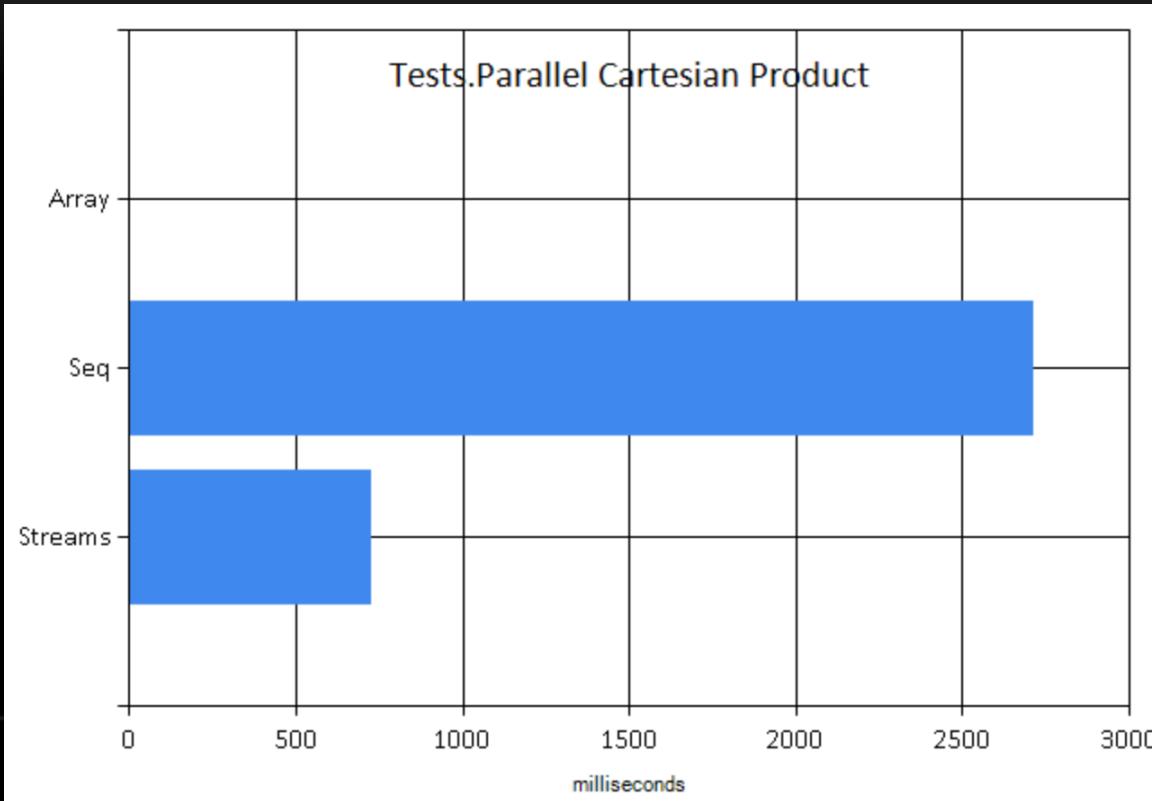




Parallel Sum of Squares



Parallel Cartesian Product





DEMO



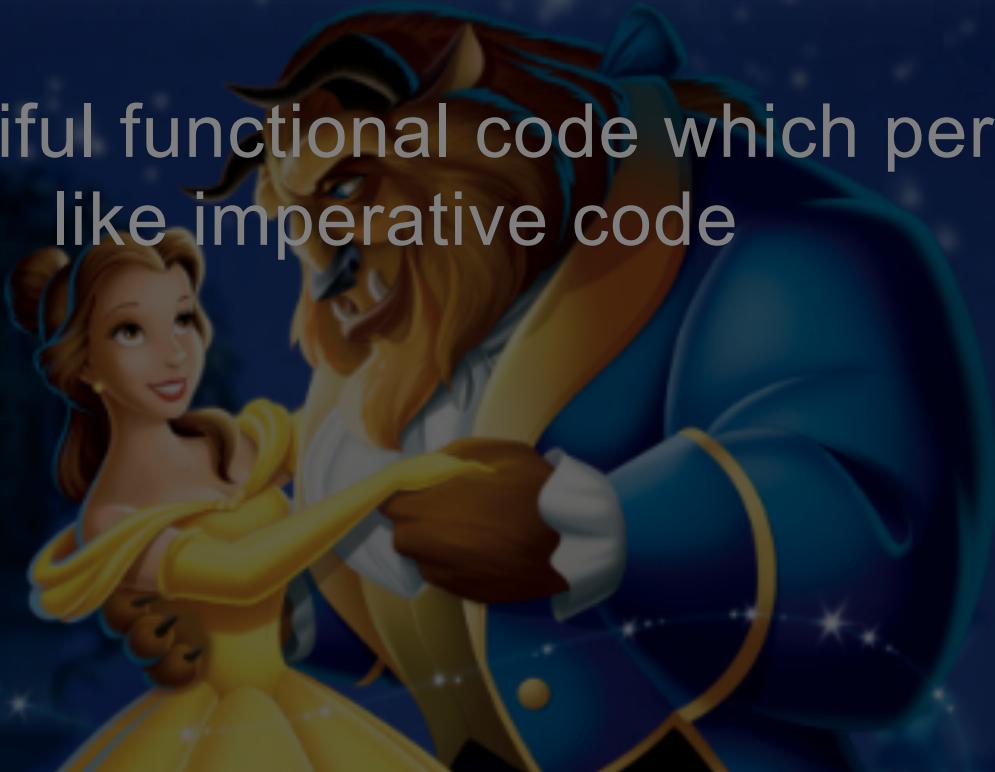
Streams are lightweight and powerful

Sequential & Parallel



Beauty and the Beast

Write beautiful functional code which performs
like imperative code





Almost...

Depends on the compiler's ability to inline.



Can we make compilers smarter?



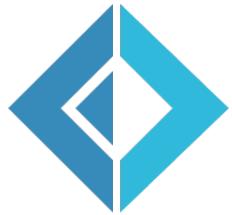


GOALS - CONCLUSION

Streams are
lightweight and
powerful

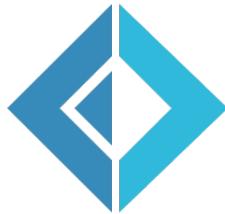
CPS is a great,
useful but miss-
used tool
- to extend functionality

There is always
room to improve
performance



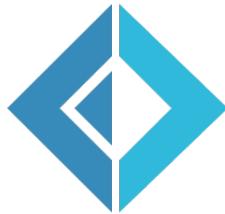
The tools we use have a profound (and devious!) influence on our thinking habits, and, therefore, on our thinking abilities.

-- Edsger Dijkstra



REFERENCES

- <https://github.com/nessos/Streams>
 - <https://github.com/biboudis/sml-streams>
 - <http://code.haskell.org/~dons/papers/icfp088-coutts.pdf>
 - <https://github.com/nessos/LinqOptimizer>
 - <http://arxiv.org/abs/1406.6631>
-



HOW TO REACH ME



github.com/rikace/Presentations/FastStreams

@DCFsharp @TRikace

rterrell@microsoft.com

That's all Folks!